

# Long-Only Equity Capacity Methodology

Distributional Liquidity Migration, Execution Horizons and Alpha Decay

Rodrigo Dupleich\*

4 October 2026

## Abstract

This document defines two complementary implementations of the methodology for estimating the capacity of a long-only equity strategy. It separates the methodology from portfolio-specific runs. Each run should produce a separate report containing the supplied inputs, fitted parameters, diagnostics, scenarios, thresholds and warnings. Appendices A–C define the user input contract and distinguish supplied, policy and fitted parameters.

## Contents

|           |                                                         |           |
|-----------|---------------------------------------------------------|-----------|
| <b>1</b>  | <b>Framework overview</b>                               | <b>3</b>  |
| <b>2</b>  | <b>Strategy economics</b>                               | <b>3</b>  |
| <b>3</b>  | <b>Distributional liquidity migration</b>               | <b>3</b>  |
| 3.1       | Burr reference model . . . . .                          | 3         |
| 3.2       | Joint ADV and parent-order calibration . . . . .        | 4         |
| 3.3       | Scenario scaling and migration . . . . .                | 5         |
| <b>4</b>  | <b>Impact model</b>                                     | <b>5</b>  |
| <b>5</b>  | <b>Execution horizon</b>                                | <b>5</b>  |
| <b>6</b>  | <b>Alpha decay</b>                                      | <b>6</b>  |
| <b>7</b>  | <b>Net capacity</b>                                     | <b>6</b>  |
| 7.1       | Burr implementation . . . . .                           | 6         |
| 7.2       | ADV-calibrated implementation . . . . .                 | 7         |
| <b>8</b>  | <b>Required run-report contents</b>                     | <b>8</b>  |
| <b>9</b>  | <b>Validation requirements</b>                          | <b>9</b>  |
| <b>10</b> | <b>Current calibration and dashboard interpretation</b> | <b>9</b>  |
| <b>11</b> | <b>Running the dashboard implementations</b>            | <b>9</b>  |
| 11.1      | Browser dashboard . . . . .                             | 10        |
| 11.2      | Streamlit dashboard . . . . .                           | 10        |
| <b>A</b>  | <b>User-supplied data inputs</b>                        | <b>10</b> |
| <b>B</b>  | <b>User-set policy and scenario parameters</b>          | <b>11</b> |

---

\*with the help of ChatGTP

**C Fitted parameters and outputs****12**

## 1 Framework overview

The framework maps assets under management (AUM) into a distribution of parent-order participation, execution horizons, implementation costs and captured alpha:

$$A \rightarrow F_A(p) \rightarrow D(p) \rightarrow \{\bar{c}(A), R_g(A)\} \rightarrow R_{\text{net}}(A), IR_{\text{net}}(A). \quad (1.1)$$

Turnover  $T$  is one-way annual turnover, so purchases and sales are approximated by  $2TA$ . In the ADV branch this is intended trading before an execution-horizon constraint. Observed cost is expressed in basis points per executed notional; the ADV scenario cost below includes completion and is expressed per intended traded dollar.

The model consumes either aggregate ADV buckets or, preferably, parent-order data. In the Burr implementation, aggregate bucket shares are shares of executed notional. The ADV-calibrated implementation jointly fits parent-order counts and traded values, using universe ADV percentiles as a liquidity prior. These are distinct weighting measures and are reported separately.

## 2 Strategy economics

At current AUM  $A_0$ , the key quantities are gross relative return  $R_g$ , tracking error  $TE$ , turnover  $T$  and average holdings  $N$ :

$$IR_g = \frac{R_g}{TE}, \quad \alpha_{\text{holding}} = \frac{R_g}{N}, \quad \alpha_{\text{cycle}} = \frac{R_g}{T}, \quad \mathcal{H} \approx \frac{1}{T}. \quad (2.1)$$

Here  $\mathcal{H}$  denotes the approximate holding period in years;  $H$  below denotes the number of holdings. Current gross alpha is treated as already reflecting the current implementation schedule. Alpha decay at alternative AUM levels is therefore measured relative to the current schedule.

## 3 Distributional liquidity migration

Capacity depends on parent-order participation  $p = Q/\text{ADV}$ , not ADV in isolation. Modelling ADV alone would require a joint model for parent-order size and its dependence on ADV.

### 3.1 Burr reference model

The aggregate implementation fits a Burr XII distribution:

$$F(p; c, d, \lambda) = 1 - \left[ 1 + \left( \frac{p}{\lambda} \right)^c \right]^{-d}. \quad (3.1)$$

It offers a flexible body and Pareto-like survival tail  $1 - F(p) \sim p^{-cd}$ . Empirical work finds power-law tails in trade-size distributions [1]. The calibration constrains  $cd \geq 1.25$ , preserving a heavy tail while ensuring finite mean participation.

The distribution changes with AUM through

$$\lambda(A) = \lambda_0 \left( \frac{A}{A_0} \right)^\eta, \quad d(A) = d_0 \left( \frac{A}{A_0} \right)^{-\kappa}. \quad (3.2)$$

The scale elasticity  $\eta$  captures how participation grows with AUM. Tail drift  $\kappa$  is reserved for structural changes in liquidity composition. Bucket shares follow smoothly from

$$w_i(A) = F_A(u_i) - F_A(\ell_i). \quad (3.3)$$

The participation distribution enters the capacity calculation through its quantile representation. The distribution is weighted by executed notional: each equal-width quantile interval represents

an equal fraction of traded value, not an equal number of trades. At each AUM, set  $F_A(p) = F(p; c, d(A), \lambda(A))$ . Its inverse turns a fixed rank  $q \in (0, 1)$  into participation:

$$p_A(q) = F_A^{-1}(q) = \lambda(A) \left[ (1 - q)^{-1/d(A)} - 1 \right]^{1/c}. \quad (3.4)$$

The same rank is used at current and alternative AUM. This is a rank-preserving scenario assumption, not an observation that the same individual orders recur. It supplies the paired participations needed for the horizon ratio in equation (5.2). When  $\kappa = 0$ ,  $p_A(q) = p_{A_0}(q)(A/A_0)^\eta$ ; with tail drift, the full inverse in equation (3.4) is required.

Bucket weights in equation (3.3) summarize this distribution for charts. The continuous model does not multiply them into its costs a second time: the notional weights are already represented by the quantile average made explicit in Section 7.

### 3.2 Joint ADV and parent-order calibration

Let  $n_b$  and  $V_b$  be the annual parent-order count and traded dollars in participation bucket  $b = (\ell_b, u_b]$ . Let  $A_\star$ ,  $H_\star$  and  $T_\star$  denote the fixed historical AUM, average holdings and one-way turnover. These historical quantities are distinct from editable scenario inputs. The sample determines

$$N_\star = \sum_b n_b, \quad V_\star = \sum_b V_b, \quad T_\star = \frac{V_\star}{2A_\star}, \quad \bar{Q}_b = \frac{V_b}{n_b}. \quad (3.5)$$

Universe ADV quantiles describe the cross-section of eligible stocks, not the frequency or value of trades. For supplied percentile fractions  $r_j$  and dollar ADVs  $L_j$ , interpolate between successive points in log dollars:

$$L(r) = \exp\{(1 - t) \log L_j + t \log L_{j+1}\}, \quad t = \frac{r - r_j}{r_{j+1} - r_j}. \quad (3.6)$$

Let  $U$  be the number of eligible stocks. The Streamlit implementation converts the continuous percentile curve into a finite stock-count grid with

$$M = \min(U, 1000), \quad r_m = \frac{m - 1/2}{M}, \quad m = 1, \dots, M. \quad (3.7)$$

Thus universes of at most 1,000 names use one midpoint rank per eligible stock; larger universes use a 1,000-point approximation. Each rank is combined with 16 evenly spaced interior participation points within each bucket. A candidate cell  $k$  in bucket  $b$  has ADV  $L_k$ , participation  $p_k \in (\ell_b, u_b)$  and ticket  $Q_k = p_k L_k$ . Its positive prior weight is proportional to

$$\pi_{bk} \propto \exp \left\{ \beta(2r_k - 1) - \frac{1}{2} \left[ \frac{\log(Q_k/\bar{Q}_b)}{1.5} \right]^2 \right\}. \quad (3.8)$$

The fixed log-ticket width 1.5 is a modelling choice. The liquidity-preference parameter  $\beta \in \{-2, 0, 2\}$  favours illiquid, neutral or liquid ranks before calibration. For each bucket, minimum-relative-entropy weights satisfy

$$\min_{w_{bk} \geq 0} \sum_k w_{bk} \log \frac{w_{bk}}{\pi_{bk}}, \quad \sum_k w_{bk} = 1, \quad \sum_k w_{bk} Q_k = \bar{Q}_b. \quad (3.9)$$

The implementation solves the one-dimensional exponential tilt

$$w_{bk}(\theta_b) = \frac{\pi_{bk} \exp(\theta_b Q_k / Q_{b,\max})}{\sum_h \pi_{bh} \exp(\theta_b Q_h / Q_{b,\max})}, \quad \nu_{bk} = n_b w_{bk}. \quad (3.10)$$

Thus  $\sum_k \nu_{bk} = n_b$  and  $\sum_k \nu_{bk} Q_k = V_b$ . Fractional counts are quadrature weights, not claims that fractional historical orders occurred. A bucket whose mean ticket lies outside the candidate support is reported as infeasible. Changing preference refits the weights. The observations do not uniquely identify the joint distribution: the prior resolves an underidentified allocation. It does not impose equal trading shares across universe deciles or reconstruct the actual 40-stock portfolio.

### 3.3 Scenario scaling and migration

For scenario AUM  $A$ , holdings  $H$  and decimal one-way turnover  $T$ , the ADV model requires  $1 \leq H \leq U$  and uses

$$s_Q = \frac{A}{A_\star} \frac{H_\star}{H}, \quad s_N = \frac{H}{H_\star} \frac{T}{T_\star}, \quad Q'_k = s_Q Q_k, \quad \nu'_k = s_N \nu_k, \quad p'_k = \frac{Q'_k}{v L_k}. \quad (3.11)$$

Here  $v > 0$  is the market-volume factor (1 for observed ADV, 0.5 for a 50% volume stress). Higher turnover increases frequency, rather than ticket size, under this convention. Higher holdings decreases tickets and increases frequency. These behavioural assumptions imply

$$\sum_k \nu'_k Q'_k = s_N s_Q V_\star = 2AT. \quad (3.12)$$

For positive turnover, count and intended-notional migration shares are

$$f_b(A) = \frac{\sum_k \nu'_k \mathbf{1}_{\{p'_k \in b\}}}{\sum_k \nu'_k}, \quad w_b(A) = \frac{\sum_k \nu'_k Q'_k \mathbf{1}_{\{p'_k \in b\}}}{\sum_k \nu'_k Q'_k}. \quad (3.13)$$

At zero turnover the interface displays zero trading shares. It continues to show hypothetical order-level execution diagnostics. An additional bucket retains participation above 100% ADV. The ADV engine does not apply Burr elasticities  $\eta$  or  $\kappa$ ; one annual sample cannot identify their response to AUM growth. Frequency changes alone leave percentage migration unchanged.

## 4 Impact model

Expected and realised costs are fitted separately using

$$c(p) = a + bp^\gamma, \quad (4.1)$$

subject to positive parameter bounds. Square-root impact is an important benchmark, while the implementation permits the exponent to be fitted [2]. At current AUM, the continuous distributional expectation is calibrated back to the observed notional-weighted bucket cost.

## 5 Execution horizon

In the Burr reference model, for target daily participation  $\rho$  and maximum horizon  $D_{\max}$ ,

$$D(p) = \min\left(D_{\max}, \max\left(1, \left\lceil \frac{p}{\rho} \right\rceil\right)\right). \quad (5.1)$$

Current costs already embed the current schedule, so counterfactual effective participation is

$$p_{\text{eff},j}(A) = p_j(A) \frac{D_j(A_0)}{D_j(A)}. \quad (5.2)$$

The rule exposes the cost-versus-delay trade-off studied in optimal execution [3]. A production version can replace it with direct optimisation.

In the ADV-calibrated model, the participation limit is enforced without treating truncated orders as complete. Required days, scheduled days and completion fraction are

$$D_k^{\text{req}} = \max(1, \lceil p'_k / \rho \rceil), \quad d_k = \min(D_{\max}, D_k^{\text{req}}), \quad e_k = \min(1, D_{\max} \rho / p'_k). \quad (5.3)$$

The implementation floors  $D_{\max}$  to a positive integer. The dashboard's average-execution statistic is the count-weighted fractional completion time

$$\bar{\tau}(A) = \frac{\sum_k \nu'_k (p'_k / \rho)}{\sum_k \nu'_k}. \quad (5.4)$$

It may therefore be below one day when most parent orders use less than one daily participation allowance. The integer  $D_k^{\text{req}}$  remains the conservative scheduling input for alpha decay, completion and multi-day classifications. Impact uses  $\tilde{p}_k = \min(p'_k, \rho)$  as a daily-rate proxy. This proxy approximates a participation-capped schedule and does not model the smaller final slice separately. Alpha capture uses equal slices over  $d_k$  for the completed fraction. These two scheduling approximations should be refined together when daily executions are available.

## 6 Alpha decay

Let  $H_\alpha$  be signal half-life in trading days and  $\delta = \log(2)/H_\alpha$ . For equal daily slices over  $D$  days, average alpha capture is

$$g(D, H_\alpha) = \frac{1}{D} \sum_{t=0}^{D-1} e^{-\delta t} = \frac{1 - e^{-\delta D}}{D(1 - e^{-\delta})}. \quad (6.1)$$

For the Burr model, normalising each notional rank to current implementation gives

$$R_g(A) = R_g(A_0) \mathbb{E} \left[ \frac{g(D_j(A), H_{\alpha,j})}{g(D_j(A_0), H_{\alpha,j})} \right]. \quad (6.2)$$

Aggregate runs currently weight alpha by executed notional. Trade-level runs should weight by each order's forecast alpha contribution. Predictors with different mean-reversion speeds imply different desired trading speeds [4].

For the ADV model, the completed alpha fraction is  $e_k g(d_k, H_\alpha)$ ; unfinished notional contributes zero captured alpha. The engine normalises the aggregate capture to its historical aggregate, rather than averaging the individual ratios used by Burr. The resulting equation is specified below.

## 7 Net capacity

### 7.1 Burr implementation

The participation distribution specified in Section 3 determines  $p_A(q)$  through equations (3.1), (3.2) and (3.4). For each notional rank  $q$ , the execution rule in equation (5.1) is evaluated at current and alternative AUM:

$$\begin{aligned} D_A(q) &= D(p_A(q)), & D_0(q) &= D(p_{A_0}(q)), \\ p_{\text{eff},A}(q) &= p_A(q) \frac{D_0(q)}{D_A(q)}. \end{aligned} \quad (7.1)$$

Thus distributional migration has two channels: it changes effective participation and therefore impact, and it changes execution days and therefore alpha capture.

Let  $\bar{c}_{\text{obs},0} = \sum_i w_{i,0}^{\text{obs}} c_i^{\text{obs}}$  be the observed current notional-weighted cost in basis points. The impact function  $c(p)$  in equation (4.1) is the raw fitted curve. Its baseline calibration multiplier is

$$K = \frac{\bar{c}_{\text{obs},0}}{\int_0^1 c(p_{A_0}(q)) dq}. \quad (7.2)$$

Expected and realised curves each use their own multiplier. Making the expectation explicit, the expected cost and captured gross alpha are

$$\bar{c}(A) = K \int_0^1 c(p_{\text{eff},A}(q)) dq, \quad (7.3)$$

$$R_g(A) = R_g(A_0) \int_0^1 \frac{g(D_A(q), H_\alpha)}{g(D_0(q), H_\alpha)} dq. \quad (7.4)$$

Here  $g$  is equation (6.1), and a common half-life is used for the aggregate model. These integrals are the notional-weighted expectations previously written as  $\mathbb{E}[\cdot]$ . Integrating over uniform ranks is equivalent to integrating over the participation distribution supplied by Section 3, with current and alternative horizons paired at the same rank.

The expectations are evaluated numerically using  $J$  equally weighted midpoint quantiles:

$$\begin{aligned} q_j &= \frac{j - \frac{1}{2}}{J}, \quad j = 1, \dots, J, & K_J &= \frac{\bar{c}_{\text{obs},0}}{J^{-1} \sum_{j=1}^J c(p_{A_0}(q_j))}, \\ \bar{c}_J(A) &= \frac{K_J}{J} \sum_{j=1}^J c(p_{\text{eff},A}(q_j)), \\ R_{g,J}(A) &= \frac{R_g(A_0)}{J} \sum_{j=1}^J \frac{g(D_A(q_j), H_\alpha)}{g(D_0(q_j), H_\alpha)}. \end{aligned} \quad (7.5)$$

The factor  $1/J$  supplies the notional weight. There is no additional bucket-share multiplier. Using the same numerical grid in the calibration denominator and cost average ensures  $\bar{c}_J(A_0) = \bar{c}_{\text{obs},0}$  and  $R_{g,J}(A_0) = R_g(A_0)$ . The continuous impact expectation requires  $c d(A) > \gamma$  for the Burr tail; the finite-mean participation constraint alone does not guarantee this for every fitted exponent or AUM. Finite-grid runs should check sensitivity to  $J$ , particularly under tail thickening.

The resulting impact estimate is converted from cost per executed notional into annual cost per unit of portfolio capital using one-way turnover:

$$C(A) = \frac{2T\bar{c}(A)}{10,000}. \quad (7.6)$$

The conversion by 10,000 changes basis points into a decimal return. Net alpha and IR are

$$R_{\text{net}}(A) = R_g(A) - C(A), \quad IR_{\text{net}}(A) = \frac{R_{\text{net}}(A)}{TE}. \quad (7.7)$$

The quantities  $R_g$ ,  $C$  and  $TE$  are expressed as decimal annual returns, and  $T$  as decimal one-way turnover. Distributional migration therefore affects both components of net return: impact in  $C(A)$  and delay in  $R_g(A)$ . Tracking error and turnover remain fixed assumptions in this specification.

For example, consider a notional rank whose participation increases from 8% to 16% ADV. With a 10% daily target and a non-binding horizon cap, its horizon increases from one to two days. Effective participation becomes  $16\% \times 1/2 = 8\%$ , so its modeled impact is unchanged, while its alpha-capture ratio becomes  $g(2, H_\alpha)/g(1, H_\alpha) < 1$ . Migration can therefore lower net IR through delay even when scheduling offsets the impact increase. This is an illustrative rank, not a claim about the entire portfolio.

A run report should show the complete capacity curve and explicitly selected thresholds rather than presenting a single universal capacity number.

## 7.2 ADV-calibrated implementation

Section 3.2 supplies calibrated synthetic counts and tickets. Section 3.3 scales them into participation, and equation (5.3) supplies completion and delay. Because every ticket scales by the same factor, the intended-notional weights are invariant:

$$\omega_k = \frac{\nu_k Q_k}{V_\star}, \quad \sum_k \omega_k = 1. \quad (7.8)$$

Evaluate the historical reference at  $A = A_\star$ ,  $H = H_\star$ ,  $v = 1$ , under the selected daily participation limit, horizon and half-life. Use a subscript  $\star$  for those reference execution quantities. Then

$$G_\star = \sum_k \omega_k e_{k,\star} g(d_{k,\star}, H_\alpha), \quad K_{\text{ADV}} = \frac{\sum_b (V_b/V_\star) c_b^{\text{obs}}}{\sum_k \omega_k e_{k,\star} c(\tilde{p}_{k,\star})}. \quad (7.9)$$

The current dashboard uses expected bucket costs for this anchor. Its captured gross alpha and cost in basis points per intended traded dollar are

$$\begin{aligned}
 R_g^{\text{ADV}}(A) &= R_{g,\star} \frac{\sum_k \omega_k e_k g(d_k, H_\alpha)}{G_\star}, \\
 \bar{c}^{\text{ADV}}(A) &= K_{\text{ADV}} \sum_k \omega_k e_k c(\tilde{p}_k), \\
 U(A) &= \sum_k \omega_k (1 - e_k).
 \end{aligned} \tag{7.10}$$

Here  $U$  is the unfinished fraction of intended notional. Substituting  $\bar{c}^{\text{ADV}}$  into equations (7.6) and (7.7) yields annual drag, net alpha and net IR. In this branch  $2AT$  is intended annual trading; completion is already included through  $e_k$ , and must not be applied again. Declining costs caused by non-completion are not evidence of improved execution:  $U$  and captured alpha must be read alongside costs.

**Interpretive note: after-delay versus net alpha.** After-delay alpha contains the loss from slower implementation but has not yet deducted market impact. Net alpha subtracts annual trading-cost drag. In the default Swiss reference case,  $R_{g,\star} = 0.5000\%$ , captured alpha is 100% by reference normalisation, expected impact is 8.69 bp and  $T = 13.9\%$ . Equation (7.6) gives 2.42 bp of annual drag, so after-delay alpha is 0.5000% and net alpha is approximately 0.4758%. Positive implementation cost therefore places the net-alpha curve below the after-delay curve even at the reference AUM.

At the historical reference the normalisation reproduces the supplied alpha and weighted cost. Changing execution-policy controls also changes  $G_\star$  and  $K_{\text{ADV}}$ . Consequently the current engine compares relative AUM scenarios under each selected policy; it is not an absolute historical execution-policy backtest. Neither aggregate branch identifies overlapping orders in one stock, time-varying volume, actual holdings selection or order-level alpha contributions.

## 8 Required run-report contents

Every portfolio run report must include:

1. run identity, timestamp, code version and input source;
2. supplied strategy, liquidity, execution and alpha-decay inputs;
3. fitted distribution and impact parameters with fit diagnostics;
4. current-AUM reconciliation to observed costs and alpha;
5. AUM scenario table and continuous capacity curves;
6. engine identity, count and dollar calibration residuals, migration and execution-horizon diagnostics;
7. for ADV runs, preference prior, volume stress, uncapped required days and unfinished notional;
8. selected capacity thresholds and exact crossing estimates;
9. data-quality, extrapolation and model-risk warnings.

## 9 Validation requirements

A valid run should confirm that notional shares sum to one, AUM and half-life are positive, participation buckets are ordered and non-overlapping, expected and realised costs share a common benchmark, and current-AUM impact reconciles to the supplied weighted cost. The report should flag missing trade-level alpha weights, large extrapolations, horizons at  $D_{\max}$ , unstable tail fits and sensitivity to user policy parameters.

## 10 Current calibration and dashboard interpretation

The supplied annual sample contains 971 parent orders and USD 944.475 million of two-way traded value at historical AUM USD 3.4 billion. It implies  $T_{\star} = 13.889338\%$  one-way turnover. The dashboard’s rounded 13.9% scenario produces USD 945.2 million of intended annual trading and approximately 971.7 expected orders at 40 holdings. This rounding difference is intentional; the calibration table itself reconciles to the exact historical totals.

The 203-stock eligible universe has a one-year ADV lookback. It therefore uses 203 midpoint stock ranks in equation (3.7). The supplied percentile boundaries are shown below; they are not mean ADVs within deciles.

| Q0    | Q10   | Q20   | Q30   | Q40   | Q50   | Q60   | Q70   | Q80   | Q90    | Q100    |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|--------|---------|
| 0.020 | 0.050 | 0.130 | 0.800 | 1.000 | 1.420 | 2.000 | 4.000 | 6.000 | 50.000 | 200.000 |

ADV in USD millions per day.

At 10% daily participation, equation (5.4) gives approximately 0.17 trading day for the calibrated current-AUM parent orders and 0.35 day at twice current AUM. These are count-weighted fractional completion times. Whole-day execution buckets still round orders up for alpha-delay and multi-day diagnostics. With 40 assumed holdings and unchanged frequency, doubling AUM doubles every synthetic ticket; annual order count remains unchanged. With the same AUM and twice the holdings, tickets halve and frequency doubles. With twice the turnover alone, frequency doubles but participation shares are unchanged. These identities are model assumptions, not estimated portfolio responses.

The main dashboard retains Burr as the default comparison. Selecting ADV-calibrated migration updates the overview and alpha-decay engine. The liquidity demo provides independent ADV and Burr controls; resetting a demo copies current dashboard assumptions without modifying the main case. The ADV demo shows count and dollar migration, historical reconciliation, execution feasibility and relative alpha capture. The impact demo illustrates  $a + bp^{\gamma}$  independently. The alpha-decay demo adds a waiting-time factor  $2^{-t/H_{\alpha}}$  to implemented gross alpha and displays 126 trading days, approximately six months; this is an explanatory extension, not an additional waiting penalty in the capacity engine.

Validation covers historical counts and values, percentile interpolation, annual trading conservation, volume-stress equivalence and infeasible calibration inputs. Quadrature density, liquidity preference and ticket-prior width require sensitivity review before operational use.

## 11 Running the dashboard implementations

The methodology is distributed through two live dashboard implementations. The TypeScript browser version and the Python Streamlit version both evaluate the model interactively from the source parameters. Both implementations use USD, the same live bucket observations, the same universe ADV quantiles and the same Burr and ADV calibration inputs.

### 11.1 Browser dashboard

The browser dashboard requires Node.js 22.13 or later. From the project root, install the locked dependencies and start the local development server with

```
cd dashboard
npm ci
npm run dev -- --host 127.0.0.1 --port 3011
```

The dashboard is then available at <http://localhost:3011>. The process remains attached to the terminal and is stopped with **Ctrl+C**. Before distributing a change, the production build and the model checks should be run from the `dashboard` directory:

```
npm run build
node --experimental-strip-types --test tests/adv-migration.test.mjs
```

The browser version recalculates immediately when the user changes AUM, holdings, turnover, participation, execution horizon, alpha half-life or the available model-specific controls.

### 11.2 Streamlit dashboard

The Streamlit implementation requires Python 3.11 or later. From the project root, create an isolated environment, install the declared dependencies and start the application with

```
cd streamlit
python3 -m venv .venv
source .venv/bin/activate
python -m pip install -r requirements.txt
streamlit run app.py
```

On Windows PowerShell, the environment is activated with `.venv\Scripts\Activate.ps1`. Streamlit normally opens <http://localhost:8501>; the server is stopped with **Ctrl+C**. The Streamlit version contains the Overview, Liquidity migration, Demo building blocks, Input data and Strategy creator areas. The built-in Swiss Strategy has 203 eligible stocks. The Strategy creator saves a named case containing dashboard defaults, liquidity buckets, universe size and ADV percentiles in the local `data/strategies.user.json` file; saving an existing name replaces its local record. The sidebar selector loads the selected case throughout the Capacity Lab. The saved universe size constrains holdings and sets the finite grid in equation (3.7). The application reads the live bucket and ADV JSON files used by the TypeScript dashboard and recalculates scenarios when AUM, turnover, holdings, participation, horizon, half-life or model controls change. The Streamlit model checks are run from its directory with

```
python -m unittest discover -p 'test_*.py'
```

## A User-supplied data inputs

Table 1: Required and optional user data

| Input                      | Requirement | Unit / format | Definition and validation                                                                            |
|----------------------------|-------------|---------------|------------------------------------------------------------------------------------------------------|
| Current AUM $A_0$          | Required    | Currency bn   | Average deployed AUM represented by the liquidity sample; must be positive.                          |
| Gross expected alpha $R_g$ | Required    | Decimal p.a.  | Expected relative return before modeled impact; document whether existing live slippage is embedded. |

| Input                           | Requirement                  | Unit / format                 | Definition and validation                                                                                     |
|---------------------------------|------------------------------|-------------------------------|---------------------------------------------------------------------------------------------------------------|
| Tracking error $TE$             | Required                     | Decimal p.a.                  | Annualised ex-ante or realised relative volatility; must be positive and methodologically aligned with alpha. |
| One-way turnover $T$            | Required                     | Decimal p.a.                  | Purchases or sales divided by average AUM. Do not supply two-way turnover without changing the convention.    |
| Average holdings $N$ or $H$     | Required for ADV scaling     | Count                         | Historical and scenario average positions; distinct from eligible universe size.                              |
| Bucket lower and upper bounds   | Required for aggregate run   | Fraction of ADV               | Ordered, non-overlapping participation intervals covering the full positive range.                            |
| Executed-notional share         | Required for aggregate run   | Fraction summing to 1         | Annual executed-notional mass in each bucket; not trade-count frequency.                                      |
| Representative participation    | Required for impact fit      | Fraction of ADV               | Prefer the notional-weighted mean participation within each bucket.                                           |
| Expected impact                 | Required                     | bp / executed notional        | Expected implementation cost using a consistent arrival-price or other declared benchmark.                    |
| Realised impact                 | Recommended                  | bp / executed notional        | Realised cost using the same benchmark as expected impact.                                                    |
| Parent-order notional or shares | Required for trade-level run | Currency / shares             | Parent order before child-order slicing, net of cancellations according to a declared convention.             |
| Security ADV                    | Required for trade-level run | Currency or shares/day        | ADV measured consistently with order units, including declared lookback and lag.                              |
| Annual bucket counts and values | Required for ADV calibration | Count / currency              | Parent orders and two-way traded dollars for the same historical period and strategy.                         |
| Universe ADV quantiles          | Required for ADV calibration | Percentile / currency per day | Include minimum and maximum, universe size and lookback; specify the daily-volume statistic.                  |
| Order timestamps                | Recommended                  | Datetime                      | Decision, release, first fill and final fill times.                                                           |
| Actual execution horizon        | Recommended                  | Trading days                  | Number of sessions over which the parent order was worked.                                                    |
| Forecast alpha                  | Recommended                  | bp or return                  | Expected contribution at the order decision timestamp.                                                        |
| Signal family                   | Recommended                  | Category                      | Used to assign or estimate signal-specific decay.                                                             |
| Observed alpha path             | Preferred                    | Return by horizon             | Forward alpha after 1, 2, 3, 5, 10 or more trading days for empirical half-life estimation.                   |
| Side, spread and volatility     | Recommended                  | Category / bp / decimal       | Used for richer asymmetric impact and execution-risk models.                                                  |

## B User-set policy and scenario parameters

Table 2: Parameters chosen or approved by the user

| Parameter                             | Example | Unit       | Interpretation                                                                                                          |
|---------------------------------------|---------|------------|-------------------------------------------------------------------------------------------------------------------------|
| Participation-scale elasticity $\eta$ | 0.85    | Elasticity | Response of the Burr scale to AUM. Set to 1 for proportional scaling; estimate from portfolio rescaling where possible. |

| Parameter                            | Example  | Unit         | Interpretation                                                                                                                 |
|--------------------------------------|----------|--------------|--------------------------------------------------------------------------------------------------------------------------------|
| Tail-thickening elasticity $\kappa$  | 0.00     | Elasticity   | Structural change in upper-tail liquidity composition. Keep at zero unless supported by portfolio simulations or observations. |
| Target daily participation $\rho$    | 10%      | ADV/day      | Execution-policy target used to determine horizon.                                                                             |
| Maximum execution horizon $D_{\max}$ | 10       | Trading days | Burr caps the horizon; ADV caps completion and reports unfinished notional.                                                    |
| Alpha half-life $H_\alpha$           | 10       | Trading days | Time for forecast alpha to halve. Prefer signal-specific empirical estimates.                                                  |
| AUM scenario grid                    | 1–100    | Currency bn  | Range over which curves and thresholds are evaluated.                                                                          |
| Liquidity preference $\beta$         | −2, 0, 2 | Prior tilt   | Illiquid, neutral or liquid preference; not an observed trading allocation.                                                    |
| Market-volume factor $v$             | 1.00     | Multiplier   | Stress scenario applied to ADV after historical calibration.                                                                   |
| Net IR threshold                     | 0.40     | Ratio        | Minimum acceptable net information ratio.                                                                                      |
| Alpha-retention threshold            | 80%      | Percentage   | Minimum fraction of current gross alpha remaining after delay and impact.                                                      |
| Impact-drag threshold                | 100      | bp p.a.      | Maximum acceptable annual return drag.                                                                                         |
| Alpha-capture threshold              | 95%      | Percentage   | Minimum gross alpha retained after execution delay, before impact.                                                             |

## C Fitted parameters and outputs

The user does not normally supply Burr parameters  $(c, d, \lambda_0)$  or impact parameters  $(a, b, \gamma)$ ; the model estimates them. Every run report should disclose them along with the tail index  $cd$ , bucket-fit residuals, impact-fit residuals, baseline calibration multiplier, threshold crossings and warnings. Parameters should be versioned with their input sample and estimation date. ADV runs additionally disclose bucket tilts, grid density, prior width, preference, count and value residuals, and the distinction between exact historical turnover and rounded scenario inputs.

## References

- [1] R. Rak et al. “Tests of scaling and universality of the distributions of trade size and share volume.” *Physical Review E* 76, 046109 (2007). doi:10.1103/PhysRevE.76.046109.
- [2] J. D. Farmer et al. “How efficiency shapes market impact.” *Quantitative Finance* 13(11), 1743–1758 (2013). arXiv:1102.5457.
- [3] R. Almgren and N. Chriss. “Optimal Execution of Portfolio Transactions.” *Journal of Risk* 3, 5–39 (2001). doi:10.21314/JOR.2001.041.
- [4] N. Gârleanu and L. H. Pedersen. “Dynamic Trading with Predictable Returns and Transaction Costs.” *Journal of Finance* 68(6), 2309–2340 (2013).